

Tutorial #2 - Explore and Visualize

October 22, 2018

1 Tutorial #2 - Explore and Visualize

Welcome to Cognitive Class Labs. This notebook is the **second** in a series of "getting started" tutorials that is designed to introduce some basic concepts and help get you familiar with using the workbench.

In this notebook, we will explore and visualize the olympic medals data that you added to your workbench in **Tutorial #1 - Get Data**. Specifically, this tutorial covers:

1. Load data in memory using a pandas DataFrame
2. Explore and manipulate data using DataFrame functions
3. Group data by columns
4. Calculate statistics over grouped data
5. Plot data using the pre-installed `matplotlib` package
6. Plot data using a third-party library

We will do this in the context of answering the following questions:

1. Which discipline and event has awarded the most gold medals?
2. Which country has won the most gold, silver, and bronze medals?

1.1 Pre-requisites

- A basic familiarity with the [Python Programming Language](#) and the [IPython Notebook](#).
- A basic understanding of [common graphical techniques](#) used for exploratory data analysis.
- A basic understanding of [matplotlib](#)

1.2 Load Data

Our first step is to load the olympic medal data into a pandas [DataFrame](#) in memory. As demonstrated in **Tutorial #1 - Get Data**, we can do this by following steps:

1. Download the olympic medal data in CSV format. Click this [Box link](#) to open the document in a new browser window.
2. Save the CSV file to your computer by clicking on the Download button.
3. Drag the CSV file from your desktop onto the workbench (Note that the CSV file appears under your **Recent Data** panel in the sidebar.)
4. Click the arrow button (>) next to the CSV file you just uploaded.
5. In the section that appears below the item, click "Rename"

6. Change the name of the file to "medals.csv" and press Enter or click outside the name.
7. Execute the code cell by clicking the () play button on the notebook toolbar, or by pressing Ctrl-Enter.

```
In [1]: import pandas
medals_df = pandas.read_csv('/resources/medals.csv')
# Prune non-data rows
medals_df = medals_df.dropna()

In [8]: medals_df.tail()

Out[8]:
```

	Year	City	Sport	Discipline	NOC	Event	Event	gender	\
2306	2006	Turin	Skiing	Snowboard	USA	Half-pipe		M	
2307	2006	Turin	Skiing	Snowboard	USA	Half-pipe		W	
2308	2006	Turin	Skiing	Snowboard	USA	Half-pipe		W	
2309	2006	Turin	Skiing	Snowboard	USA	Snowboard Cross		M	
2310	2006	Turin	Skiing	Snowboard	USA	Snowboard Cross		W	

	Medal
2306	Silver
2307	Gold
2308	Silver
2309	Gold
2310	Silver

```
In [9]: medals_df.head(6)

Out[9]:
```

	Year	City	Sport	Discipline	NOC	Event	\
0	1924	Chamonix	Skating	Figure skating	AUT	individual	
1	1924	Chamonix	Skating	Figure skating	AUT	individual	
2	1924	Chamonix	Skating	Figure skating	AUT	pairs	
3	1924	Chamonix	Bobsleigh	Bobsleigh	BEL	four-man	
4	1924	Chamonix	Ice Hockey	Ice Hockey	CAN	ice hockey	
5	1924	Chamonix	Biathlon	Biathlon	FIN	military patrol	

	Event	gender	Medal
0		M	Silver
1		W	Gold
2		X	Gold
3		M	Bronze
4		M	Gold
5		M	Silver

1.3 Explore

So what does the olympic medal data look like? We can peek at the data and its structure by looking at the first few rows. The DataFrame's `head()` method exists for this purpose.

```
In [10]: medals_df.head()
```

```

Out[10]:
   Year  City      Sport  Discipline  NOC  Event  Event  gender  \
0  1924  Chamonix  Skating  Figure skating  AUT  individual  M
1  1924  Chamonix  Skating  Figure skating  AUT  individual  W
2  1924  Chamonix  Skating  Figure skating  AUT  pairs      X
3  1924  Chamonix  Bobsleigh  Bobsleigh  BEL  four-man  M
4  1924  Chamonix  Ice Hockey  Ice Hockey  CAN  ice hockey  M

   Medal
0  Silver
1   Gold
2   Gold
3  Bronze
4   Gold

```

Each row provides:

- The year the medal was awarded
- The city where the games took place
- The sport
- The discipline
- The nationality of the medal winner
- The specific event
- The gender of the medal winner
- The type of medal

A first logical question might be, for what time period does the data apply? To find out, we turn our attention to the 'Year' column, which we can access directly from the DataFrame object. We can invoke the built-in `min()` and `max()` functions on the column (which is a pandas [Series](#)).

```
In [11]: medals_df.Year.min(), medals_df.Year.max()
```

```
Out[11]: ('1924', '2006')
```

To be more precise, we can list all years for which we have medal data.

```
In [14]: medals_df.Year.unique()
```

```
Out[14]: array(['1924', '1928', '1932', '1936', '1948', '1952', '1956', '1960',
                '1964', '1968', '1972', '1976', '1980', '1984', '1988', '1992',
                '1994', '1998', '2002', '2006'], dtype=object)
```

OK. So it looks like we have medal data for all winter olympics from 1924 through 2006.

Next question: How many medals have been awarded? The easiest way to find out is to count the rows using Python's built-in `len()` function, which returns the number of rows.

```
In [15]: len(medals_df)
```

```
Out[15]: 2311
```

ow about a sanity check? What are the distinct medal colors for all medals awarded? And how many of each color were awarded? We can answer these questions without much effort. Here we access the DataFrame's Medal column and invoke built-in functions on the resulting Series.

```
In [16]: medals_df.Medal.unique()
```

```
Out[16]: array(['Silver', 'Gold', 'Bronze'], dtype=object)
```

```
In [17]: medals_df.Medal.value_counts()
```

```
Out[17]: Gold      774
         Silver    773
         Bronze    764
         Name: Medal, dtype: int64
```

Well, these results make sense.

Let us now list all the winter olympic sports and disciplines that have awarded medals. To do this, we first use the DataFrame `groupby()` function to group all medal data by sport and discipline. We then list the keys used to identify each group.

```
In [18]: disciplines = medals_df.groupby(['Sport', 'Discipline'])
```

```
In [22]: disciplines.groups.keys()
```

```
Out[22]: dict_keys([('Biathlon', 'Biathlon'), ('Bobsleigh', 'Bobsleigh'), ('Bobsleigh', 'Skeleto
```

1.4 Question 1: Which discipline and event has awarded the most gold medals?

The first step we must take to answer this question is to filter our data. We want to ensure we only consider gold medals. We can apply a filter to our DataFrame by selecting records using a boolean indicator. We store the result in a new DataFrame.

```
In [23]: gold_df = medals_df[medals_df.Medal == 'Gold']
```

```
In [24]: gold_df.head()
```

```
Out[24]:
```

	Year	City	Sport	Discipline	NOC	Event	Event	gender	\
1	1924	Chamonix	Skating	Figure skating	AUT	individual		W	
2	1924	Chamonix	Skating	Figure skating	AUT	pairs		X	
4	1924	Chamonix	Ice Hockey	Ice Hockey	CAN	ice hockey		M	
7	1924	Chamonix	Skating	Speed skating	FIN	10000m		M	
9	1924	Chamonix	Skating	Speed skating	FIN	1500m		M	

```

Medal
1  Gold
2  Gold
4  Gold
7  Gold
9  Gold

```

```
In [25]: gold_df.tail()
```

```
Out[25]:
```

	Year	City	Sport	Discipline	NOC	Event	Event	gender	\
2301	2006	Turin	Skiing	Alpine Skiing	USA	Alpine combined		M	
2302	2006	Turin	Skiing	Alpine Skiing	USA	giant slalom		W	
2305	2006	Turin	Skiing	Snowboard	USA	Half-pipe		M	
2307	2006	Turin	Skiing	Snowboard	USA	Half-pipe		W	
2309	2006	Turin	Skiing	Snowboard	USA	Snowboard Cross		M	

	Medal
2301	Gold
2302	Gold
2305	Gold
2307	Gold
2309	Gold

Next, we want to group the gold medals by the Discipline and Event columns. The result gives us the gold medals awarded by event.

```
In [26]: by_event = gold_df.groupby(['Discipline', 'Event'])
```

We can get a glimpse of the grouped data using the DataFrame head() function.

```
In [29]: by_event.head()
```

```
Out[29]:
```

	Year	City	Sport	Discipline	NOC	\
1	1924	Chamonix	Skating	Figure skating	AUT	
2	1924	Chamonix	Skating	Figure skating	AUT	
4	1924	Chamonix	Ice Hockey	Ice Hockey	CAN	
7	1924	Chamonix	Skating	Speed skating	FIN	
9	1924	Chamonix	Skating	Speed skating	FIN	
10	1924	Chamonix	Skating	Speed skating	FIN	
14	1924	Chamonix	Skating	Speed skating	FIN	
20	1924	Chamonix	Curling	Curling	GBR	
30	1924	Chamonix	Skiing	Cross Country S	NOR	
33	1924	Chamonix	Skiing	Cross Country S	NOR	
36	1924	Chamonix	Skiing	Nordic Combined	NOR	
38	1924	Chamonix	Skiing	Ski Jumping	NOR	
40	1924	Chamonix	Biathlon	Biathlon	SUI	
41	1924	Chamonix	Bobsleigh	Bobsleigh	SUI	
44	1924	Chamonix	Skating	Figure skating	SWE	
47	1924	Chamonix	Skating	Speed skating	USA	
54	1928	St. Moritz	Ice Hockey	Ice Hockey	CAN	
55	1928	St. Moritz	Skating	Speed skating	FIN	
58	1928	St. Moritz	Skating	Speed skating	FIN	
59	1928	St. Moritz	Skating	Figure skating	FRA	
62	1928	St. Moritz	Skating	Figure skating	NOR	
66	1928	St. Moritz	Skating	Speed skating	NOR	
68	1928	St. Moritz	Skating	Speed skating	NOR	

70	1928	St. Moritz	Skiing	Cross Country S	NOR
73	1928	St. Moritz	Skiing	Nordic Combined	NOR
75	1928	St. Moritz	Skiing	Ski Jumping	NOR
79	1928	St. Moritz	Skating	Figure skating	SWE
81	1928	St. Moritz	Skiing	Cross Country S	SWE
84	1928	St. Moritz	Bobsleigh	Bobsleigh	USA
86	1928	St. Moritz	Bobsleigh	Skeleton	USA
...
2078	2006	Turin	Skiing	Nordic Combined	AUT
2087	2006	Turin	Bobsleigh	Skeleton	CAN
2106	2006	Turin	Skiing	Cross Country S	CAN
2122	2006	Turin	Skiing	Alpine Skiing	CRO
2126	2006	Turin	Skiing	Cross Country S	CZE
2130	2006	Turin	Skiing	Cross Country S	EST
2140	2006	Turin	Biathlon	Biathlon	FRA
2143	2006	Turin	Biathlon	Biathlon	FRA
2151	2006	Turin	Biathlon	Biathlon	GER
2157	2006	Turin	Biathlon	Biathlon	GER
2170	2006	Turin	Skating	Speed skating	GER
2175	2006	Turin	Skiing	Nordic Combined	GER
2185	2006	Turin	Skating	Speed skating	ITA
2194	2006	Turin	Skating	Short Track S.	KOR
2196	2006	Turin	Skating	Short Track S.	KOR
2198	2006	Turin	Skating	Short Track S.	KOR
2199	2006	Turin	Skating	Short Track S.	KOR
2235	2006	Turin	Biathlon	Biathlon	RUS
2236	2006	Turin	Biathlon	Biathlon	RUS
2250	2006	Turin	Skiing	Cross Country S	RUS
2265	2006	Turin	Skiing	Snowboard	SUI
2267	2006	Turin	Skiing	Snowboard	SUI
2268	2006	Turin	Skiing	Snowboard	SUI
2270	2006	Turin	Biathlon	Biathlon	SWE
2281	2006	Turin	Skiing	Cross Country S	SWE
2282	2006	Turin	Skiing	Cross Country S	SWE
2283	2006	Turin	Skiing	Cross Country S	SWE
2301	2006	Turin	Skiing	Alpine Skiing	USA
2305	2006	Turin	Skiing	Snowboard	USA
2309	2006	Turin	Skiing	Snowboard	USA

	Event	Event	gender	Medal
1	individual		W	Gold
2	pairs		X	Gold
4	ice hockey		M	Gold
7	10000m		M	Gold
9	1500m		M	Gold
10	5000m		M	Gold
14	combined (4 events)		M	Gold
20	curling		M	Gold

30	18km	M	Gold
33	50km	M	Gold
36	individual	M	Gold
38	K90 individual (70m)	M	Gold
40	military patrol	M	Gold
41	four-man	M	Gold
44	individual	M	Gold
47	500m	M	Gold
54	ice hockey	M	Gold
55	1500m	M	Gold
58	500m	M	Gold
59	pairs	X	Gold
62	individual	W	Gold
66	5000m	M	Gold
68	500m	M	Gold
70	18km	M	Gold
73	individual	M	Gold
75	K90 individual (70m)	M	Gold
79	individual	M	Gold
81	50km	M	Gold
84	five-man	M	Gold
86	individual	M	Gold
...
2078	Individual sprint	M	Gold
2087	individual	M	Gold
2106	sprint 1.5km	W	Gold
2122	Alpine combined	W	Gold
2126	30km	W	Gold
2130	Combined 7.5 + 7.5km mass start	W	Gold
2140	12.5km pursuit	M	Gold
2143	7.5km	W	Gold
2151	10km pursuit	W	Gold
2157	15km mass start	M	Gold
2170	Team pursuit	W	Gold
2175	Individual	M	Gold
2185	Team pursuit	M	Gold
2194	1500m	M	Gold
2196	1500m	W	Gold
2198	3000m relay	W	Gold
2199	5000m relay	M	Gold
2235	15km	W	Gold
2236	4x6km relay	W	Gold
2250	Combined 15 + 15km mass start	M	Gold
2265	Giant parallel slalom	M	Gold
2267	Giant parallel slalom	W	Gold
2268	Snowboard Cross	W	Gold
2270	12,5km mass start	W	Gold
2281	Sprint 1,5km	M	Gold

2282	Team sprint	M	Gold
2283	Team sprint	W	Gold
2301	Alpine combined	M	Gold
2305	Half-pipe	M	Gold
2309	Snowboard Cross	M	Gold

[296 rows x 8 columns]

We can easily tally the number of gold medals per event using the pandas DataFrame count() function.

```
In [30]: golds_by_event = by_event.Medal.count()
```

```
In [31]: print (golds_by_event)
```

Discipline	Event	
Alpine Skiing	Alpine combined	2
	alpine combined	14
	downhill	32
	giant slalom	30
	slalom	32
	super-G	12
Biathlon	10km	8
	10km pursuit	2
	12,5km mass start	1
	12.5km pursuit	2
	15km	5
	15km mass start	1
	20km	13
	3x7.5km relay	1
	4x6km relay	1
	4x7.5km relay	14
	7.5km	5
Bobsleigh	military patrol	1
	five-man	1
	four-man	18
	two-man	20
Cross Country S	10km	12
	10km pursuit	5
	15km	11
	15km mass start	4
	18km	6
	20km	2
	30km	5
	30km mass start	13
	3x5km relay	5
	..	
Freestyle Ski.	moguls	10
Ice Hockey	ice hockey	23

Luge	doubles	13
	singles	24
Nordic Combined	Individual	1
	Individual sprint	1
	Team	6
	individual	19
	sprint	1
Short Track S.	1000m	9
	1500m	4
	3000m relay	5
	5000m relay	5
	500m	9
Skeleton	individual	6
Ski Jumping	K120 individual (90m)	12
	K120 team (90m)	6
	K90 individual (70m)	20
Snowboard	Giant parallel slalom	4
	Half-pipe	6
	Snowboard Cross	2
	giant-slalom	2
Speed skating	10000m	19
	1000m	22
	1500m	35
	3000m	13
	5000m	26
	500m	34
	Team pursuit	2
	combined (4 events)	1

Name: Medal, Length: 77, dtype: int64

Finally, we sort the results. Here are the 10 events with the most gold medals awarded.

```
In [34]: golds_by_event.sort_values(ascending=False)
         golds_by_event.head(10)
```

```
Out[34]: Discipline Event
         Alpine Skiing Alpine combined      2
                   alpine combined     14
                   downhill           32
                   giant slalom       30
                   slalom             32
                   super-G           12
         Biathlon     10km              8
                   10km pursuit        2
                   12,5km mass start    1
                   12.5km pursuit      2

         Name: Medal, dtype: int64
```

Answer: Individual figure skating has awarded the most gold medals to olympians with 40.

1.5 Question 2: Which country has won the most gold, silver, and bronze medals?

For this question, we need to group and count medals awarded by country code. Because we need sub-totals by medal color, we must group the data by both the NOC and Medal columns. We calculate the medal counts for each group using the resulting DataFrame's `size()` function, which gives us the number of rows in each group.

```
In [35]: medals_by_country = medals_df.groupby(['NOC', 'Medal']).size()
```

```
In [36]: print (medals_by_country)
```

NOC	Medal	
AUS	Bronze	3
	Gold	3
AUT	Bronze	70
	Gold	51
	Silver	64
BEL	Bronze	3
	Gold	1
	Silver	1
BLR	Bronze	3
	Silver	3
BUL	Bronze	3
	Gold	1
	Silver	2
CAN	Bronze	43
	Gold	38
	Silver	38
CHN	Bronze	13
	Gold	4
	Silver	16
CRO	Gold	4
	Silver	3
CZE	Bronze	2
	Gold	3
	Silver	5
DEN	Silver	1
ESP	Bronze	1
	Gold	1
EST	Bronze	1
	Gold	4
	Silver	1
		..
POL	Silver	3
PRK	Bronze	1
	Silver	1
ROU	Bronze	1
RUS	Bronze	19
	Gold	33

```

    Silver    24
SLO Bronze    4
SUI Bronze   43
    Gold    38
    Silver   37
SVK Silver    1
SWE Bronze   44
    Gold    43
    Silver   31
TCH Bronze   15
    Gold     2
    Silver    8
UKR Bronze    3
    Gold     1
    Silver    1
URS Bronze   59
    Gold    78
    Silver   57
USA Bronze   58
    Gold    78
    Silver   80
UZB Gold      1
YUG Bronze    1
    Silver    3
Length: 112, dtype: int64

```

The result is a pandas Series object containing the medal counts by country.

```
In [37]: medals_by_country.head(10)
```

```

Out[37]: NOC  Medal
        AUS  Bronze    3
           Gold     3
        AUT  Bronze   70
           Gold    51
           Silver   64
        BEL  Bronze    3
           Gold     1
           Silver    1
        BLR  Bronze    3
           Silver    3
dtype: int64

```

We want to convert this Series to a DataFrame containing a column for each medal color. Fortunately, this is easy to do using the Series' `unstack()` function, which pivots the data for us by creating a column for each medal color.

```
In [38]: medals_by_country_df = medals_by_country.unstack()
        medals_by_country_df.head()
```

```
Out[38]: Medal  Bronze  Gold  Silver
        NOC
        AUS      3.0   3.0    NaN
        AUT     70.0  51.0   64.0
        BEL      3.0   1.0    1.0
        BLR      3.0   NaN    3.0
        BUL      3.0   1.0    2.0
```

Many countries do not have medals of every color, so we replace any missing data with a zero value.

```
In [39]: medals_by_country_df.fillna(0, inplace=True)
```

```
In [40]: medals_by_country_df.head()
```

```
Out[40]: Medal  Bronze  Gold  Silver
        NOC
        AUS      3.0   3.0    0.0
        AUT     70.0  51.0   64.0
        BEL      3.0   1.0    1.0
        BLR      3.0   0.0    3.0
        BUL      3.0   1.0    2.0
```

Now we can answer our question using yet another DataFrame function, `idxmax()`, which gives us the index (in this case, the country code) corresponding to the maximum count for each medal color (column in our DataFrame).

```
In [41]: medals_by_country_df.idxmax()
```

```
Out[41]: Medal
        Bronze  NOR
        Gold    NOR
        Silver  NOR
        dtype: object
```

Norway (NOR) appears to be the winner.

1.5.1 Plot

We can use the popular `matplotlib` package to produce a plot of the results.

Note: Cognitive Class Labs pre-installs many third-party Python libraries and packages. To see a list of these packages, run `!pip freeze` in a code cell.

First, tell the notebook server to render charts inline:

```
In [42]: !pip freeze
```

```
abs1-py==0.5.0
alabaster==0.7.10
anaconda-client==1.6.14
anaconda-navigator==1.8.7
```

anaconda-project==0.8.2
appdirs==1.4.3
asn1crypto==0.24.0
astor==0.7.1
astroid==1.6.3
astropy==3.0.2
attrs==18.1.0
autobahn==18.9.2
Automat==0.7.0
Babel==2.5.3
backcall==0.1.0
backports.shutil-get-terminal-size==1.0.0
basemap==1.1.0
beautifulsoup4==4.6.0
bitarray==0.8.1
bkcharts==0.2
blaze==0.11.3
bleach==1.5.0
bokeh==0.12.16
boto==2.48.0
boto3==1.7.12
botocore==1.10.84
Bottleneck==1.2.1
certifi==2018.8.24
cffi==1.11.5
chardet==3.0.4
click==6.7
cloudpickle==0.5.3
clyent==1.2.2
colorama==0.3.9
conda==4.5.11
conda-build==3.10.5
conda-verify==2.0.0
constantly==15.1.0
contextlib2==0.5.5
cryptography==2.2.2
cycller==0.10.0
Cython==0.28.2
cytoolz==0.9.0.1
dask==0.17.5
datashape==0.5.4
decorator==4.3.0
distributed==1.21.8
docutils==0.14
entrypoints==0.2.3
et-xmlfile==1.0.1
fastcache==1.0.2
filelock==3.0.4

Flask==1.0.2
Flask-Cors==3.0.4
future==0.16.0
gast==0.2.0
gevent==1.3.0
gitdb2==2.0.4
GitPython==2.1.11
glob2==0.6
gmpy2==2.0.8
greenlet==0.4.13
grpcio==1.14.1
h5py==2.8.0
heapdict==1.0.0
html5lib==0.9999999
hyperlink==18.0.0
ibm-cos-sdk==2.1.1
ibm-cos-sdk-core==2.3.0
ibm-cos-sdk-s3transfer==2.3.0
ibm-db==2.0.8a0
ibm-db-sa==0.3.3
idna==2.6
imageio==2.3.0
imagesize==1.0.0
incremental==17.5.0
ipykernel==4.8.2
ipython==6.4.0
ipython-genutils==0.2.0
ipython-sql==0.3.9
ipywidgets==7.2.1
isort==4.3.4
itsdangerous==0.24
jdcal==1.4
jedi==0.12.0
Jinja2==2.10
jmespath==0.9.3
jsonschema==2.6.0
jupyter==1.0.0
jupyter-client==5.2.3
jupyter-console==5.2.0
jupyter-core==4.4.0
jupyterlab==0.34.7
jupyterlab-cognos-dashboard-embedded==0.1.0
jupyterlab-github==0.6.1
jupyterlab-launcher==0.13.1
jupyterlab-tutorials==0.2.0
Keras==2.1.5
kiwisolver==1.0.1
lazy-object-proxy==1.3.1

llvmlite==0.23.1
locket==0.2.0
lxml==4.2.1
Mako==1.0.7
Markdown==2.6.11
MarkupSafe==1.0
matplotlib==2.2.2
mccabe==0.6.1
mistune==0.8.3
mkl-fft==1.0.0
mkl-random==1.0.1
more-itertools==4.1.0
mpmath==1.0.0
msgpack-python==0.5.6
multipledispatch==0.5.0
navigator-updater==0.2.1
nbconvert==5.3.1
nbformat==4.4.0
networkx==2.1
nltk==3.3
nose==1.3.7
notebook==5.5.0
numba==0.38.0
numexpr==2.6.5
numpy==1.14.3
numpydoc==0.8.0
odo==0.5.1
olefile==0.45.1
openpyxl==2.5.3
packaging==17.1
pandas==0.23.0
pandocfilters==1.4.2
parso==0.2.0
partd==0.3.8
path.py==11.0.1
pathlib2==2.3.2
patsy==0.5.0
pep8==1.7.1
pexpect==4.5.0
pickleshare==0.7.4
Pillow==5.2.0
pkginfo==1.4.2
pluggy==0.6.0
ply==3.11
prettytable==0.7.2
prompt-toolkit==1.0.15
protobuf==3.5.2
psutil==5.4.5

ptyprocess==0.5.2
py==1.5.3
pyarrow==0.7.1
pyasn1==0.4.4
pyasn1-modules==0.2.2
pycodestyle==2.4.0
pycosat==0.6.3
pycparser==2.18
pycrypto==2.6.1
pycurl==7.43.0.1
pydotplus==2.0.2
pyflakes==1.6.0
Pygments==2.2.0
pygpu==0.7.6
PyHamcrest==1.9.0
pylint==1.8.4
pyodbc==4.0.23
pyOpenSSL==18.0.0
pyparsing==2.2.0
pyproj==1.9.5.1
pyshp==1.2.12
PySocks==1.6.8
pytest==3.5.1
pytest-arraydiff==0.2
pytest-astropy==0.3.0
pytest-doctestplus==0.1.3
pytest-openfiles==0.3.0
pytest-remotedata==0.2.1
python-dateutil==2.7.3
pytz==2018.4
PyWavelets==0.5.2
PyYAML==3.12
pyzmq==17.0.0
QtAwesome==0.4.4
qtconsole==4.3.1
QtPy==1.4.1
quilt==2.9.4
raven==6.9.0
requests==2.18.4
rope==0.10.7
ruamel-yaml==0.15.35
s3transfer==0.1.13
scikit-image==0.13.1
scikit-learn==0.19.1
scipy==1.1.0
seaborn==0.8.1
Send2Trash==1.5.0
service-identity==17.0.0


```
simplegeneric==0.8.1
singledispatch==3.4.0.3
six==1.11.0
smmap2==2.0.4
snowballstemmer==1.2.1
sortedcollections==0.6.1
sortedcontainers==1.5.10
Sphinx==1.7.4
sphinxcontrib-websupport==1.0.1
spyder==3.2.8
SQLAlchemy==1.2.7
sqlparse==0.2.4
statsmodels==0.9.0
sympy==1.1.1
tables==3.4.3
tblib==1.3.2
tensorboard==1.8.0
tensorflow==1.8.0
termcolor==1.1.0
terminado==0.8.1
testpath==0.3.1
Theano==1.0.3
toolz==0.9.0
torch==0.4.1
torchvision==0.2.1
tornado==5.0.2
tqdm==4.26.0
traitlets==4.3.2
Twisted==18.7.0
txaio==18.8.1
typing==3.6.4
unicodcsv==0.14.1
urllib3==1.22
watson-developer-cloud==1.4.1
wcwidth==0.1.7
webencodings==0.5.1
Werkzeug==0.14.1
widgetsnbextension==3.2.1
wrapt==1.10.11
xlrd==1.1.0
XlsxWriter==1.0.4
xlwt==1.3.0
zict==0.1.3
zope.interface==4.5.0
```

telling the notebook server to render charts inline:

```
In [43]: %matplotlib inline
```


We sort the results by highest gold medal count.

```
In [44]: medals_by_country_df.sort_values('Gold', ascending=False, inplace=True)
```

Now we can use the DataFrame `plot()` function to produce our plot. We plot individual medal counts for the top 15 countries.

```
In [45]: medals_by_country_df[['Gold', 'Silver', 'Bronze'][:15]\
      .plot(kind='bar', figsize=(12,10))
```

```
Out[45]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff63e6d3f98>
```



The plot confirms that Norway has been on every step of the winter olympic podium more than any other country.

Let's take a look at another plot: the olympic medals won by Norway over time. To do this, we must filter our original medal data (containing all medal colors) to just those won by Norway, and group by Year. We then count the number of medals for each year. Amazingly, we can do all this in a single line of code.

```
In [46]: nor_medals_year = medals_df[medals_df.NOC == 'NOR'].groupby('Year').size()
```

```
In [47]: print (nor_medals_year)
```

```
Year
1924    17
1928    15
1932    10
1936    15
1948    10
1952    16
1956     4
1960     6
1964    15
1968    14
1972    12
1976     7
1980    10
1984     9
1988     5
1992    20
1994    26
1998    25
2002    25
2006    19
dtype: int64
```

```
In [48]: nor_medals_year
```


```
Out[48]: Year
1924    17
1928    15
1932    10
1936    15
1948    10
1952    16
1956     4
1960     6
1964    15
1968    14
1972    12
1976     7
1980    10
1984     9
1988     5
1992    20
1994    26
1998    25
2002    25
```

```
2006    19
dtype: int64
```

We can plot the resulting pandas Series using its built-in `plot()` function, which is just a convenience function that wraps `matplotlib`.

```
In [49]: nor_medals_year.plot(kind='bar', figsize=(12,6))
```

```
Out[49]: <matplotlib.axes._subplots.AxesSubplot at 0x7ff63e6d36d8>
```



1.6 Third-Party Visualization Libraries

It is relatively easy to create plots quickly using `pandas` and `matplotlib`. If you want to take your plots to the next level in terms of enhanced presentation, it is possible to install third-party libraries to help you.

One such library is [prettyplotlib](#), which enhances `matplotlib` plots with better default colors, etc. This can be extremely valuable if your notebook is to be used for presentations.

Note: You have the ability to install Python packages in your Cognitive Class Labs environment. Just use the `pip` package installer.

Here we use the shorthand [IPython cell magic](#) (!) to invoke `pip` to install `prettyplotlib`.

```
In [50]: # Install third party color palette
!pip install prettyplotlib
```

```
Collecting prettyplotlib
```

```
  Downloading https://files.pythonhosted.org/packages/f2/89/35079781fe5f8c4e5258b88bb0a80d4a2028
```

```
    100% || 706kB 18.5MB/s
```

```
Requirement already satisfied: matplotlib>=1.2.1 in /home/jupyterlab/conda/lib/python3.6/site-pa
```

```

Collecting brewer2mpl>=1.3.1 (from prettyplotlib)
  Downloading https://files.pythonhosted.org/packages/84/57/00c45a199719e617db0875181134fcb3aeeff
Requirement already satisfied: numpy>=1.7.1 in /home/jupyterlab/conda/lib/python3.6/site-package
Requirement already satisfied: cyclor>=0.10 in /home/jupyterlab/conda/lib/python3.6/site-package
Requirement already satisfied: pyparsing!=2.0.4,!2.1.2,!2.1.6,>=2.0.1 in /home/jupyterlab/cond
Requirement already satisfied: python-dateutil>=2.1 in /home/jupyterlab/conda/lib/python3.6/site
Requirement already satisfied: pytz in /home/jupyterlab/conda/lib/python3.6/site-packages (from
Requirement already satisfied: six>=1.10 in /home/jupyterlab/conda/lib/python3.6/site-packages (
Requirement already satisfied: kiwisolver>=1.0.1 in /home/jupyterlab/conda/lib/python3.6/site-pa
Requirement already satisfied: setuptools in /home/jupyterlab/conda/lib/python3.6/site-packages
Building wheels for collected packages: prettyplotlib
  Running setup.py bdist_wheel for prettyplotlib ... done
  Stored in directory: /home/jupyterlab/.cache/pip/wheels/76/ad/45/9fcfb9e97ecccc850d8b2fb20b8d6
Successfully built prettyplotlib
distributed 1.21.8 requires msgpack, which is not installed.
Installing collected packages: brewer2mpl, prettyplotlib
Successfully installed brewer2mpl-1.4.1 prettyplotlib-0.1.7

```

Now we create a color palette...

```

In [51]: # Use color palette "Set3" from http://bl.ocks.org/mbostock/5577023
import brewer2mpl
set3 = brewer2mpl.get_map('Set3', 'qualitative', 10).mpl_colors


```

...and use it to generate a pretty pie chart showing the percent of all medals awarded to the top 10 countries.

```

In [52]: # Plot the top 10 countries based on podium appearances.
import pylab as plt
# Calculate the total medals for each country
t = medals_by_country_df.sum(axis=1)
# Limit to top 10 medal-winning countries
t.sort_values(ascending=False)
awards = t[:10]
countries = awards.index.values
total = float(t.sum())
# Create a pie chart
pct = lambda x: '{p:1.1f}%'.format(p=(x*sum(awards)/100)/total*100)
plt.pie(awards, labels=countries, shadow=True,
        autopct=pct, colors=set3, startangle=90)
# Set aspect ratio to be equal so that pie is drawn as a circle.
plt.axis('equal')
plt.show()

```



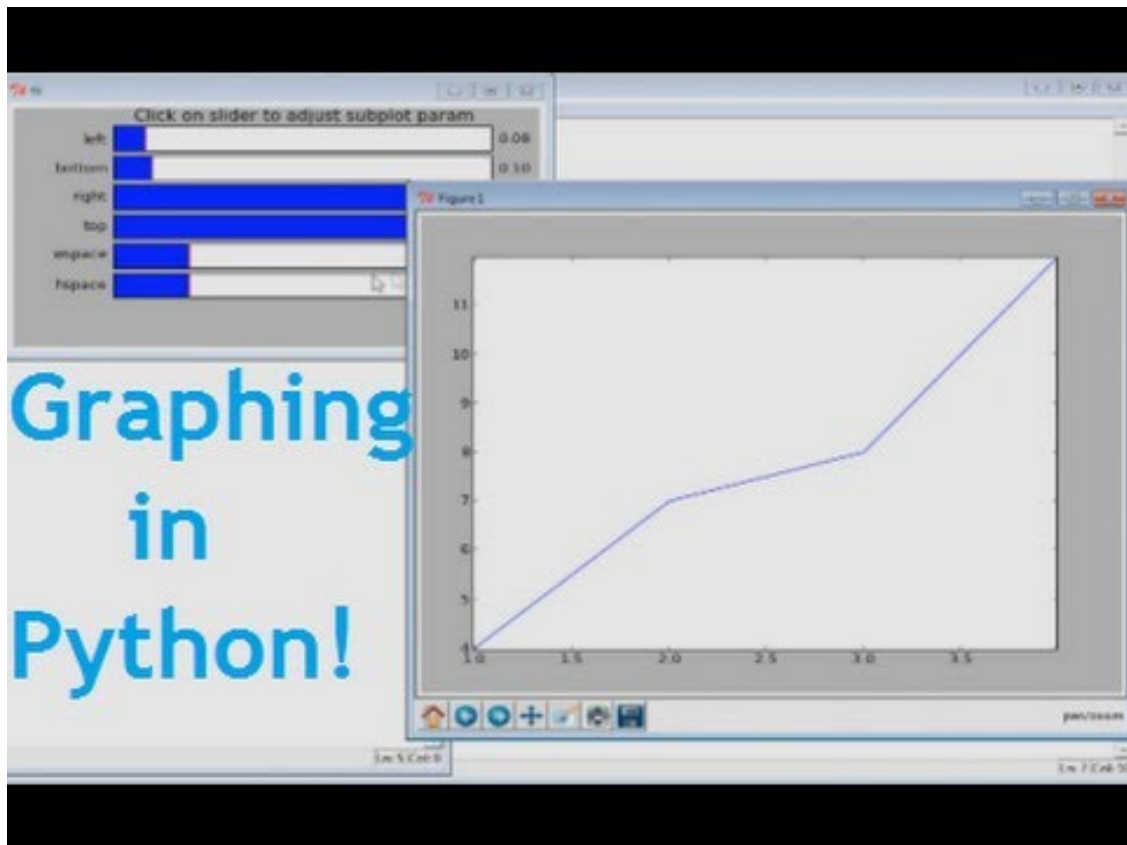
1.7 Online Tutorials

If you need further instruction on how to create visualizations in IPython notebook, there is content available online. Here is an example. The video below represents a resource for visual programming using matplotlib.

Note: The code cell below shows one way you can embed a video in your notebook.

```
In [53]: from IPython.display import YouTubeVideo
          # Matplotlib Python Tutorial Part 1: Basics and your first Graph!
          # Tutorial series by Sentdex - http://sentdex.com/about-us/
          YouTubeVideo('wAwQ-noyB98')
```

Out[53]:



1.8 Next: Organize

Our next tutorial topic will focus on how Cognitive Class Labs can help you organize your work. Visit the [Welcome](#) page to download **Tutorial #3 - Organize**.